

UMGS: UAV Multispectral 3D Gaussian Splatting for Spectral-Index Synthesis – Supplementary Material

ANONYMOUS AUTHOR(S)
SUBMISSION ID: TBD

CCS Concepts: • **Computing methodologies** → **Image-based rendering**; **Reconstruction**.

Additional Key Words and Phrases: multispectral imaging, gaussian splatting, 3D reconstruction, UAV

ACM Reference Format:

Anonymous Author(s). 2026. UMGS: UAV Multispectral 3D Gaussian Splatting for Spectral-Index Synthesis – Supplementary Material. 1, 1 (May 2026), 3 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Protocol Notes and Scope

This supplement aligns with the current main-paper draft and records additional tables that are useful for reproducibility and reviewer inspection.

The main paper focuses on a 10-scene UAV-focused benchmark: 7 raw UAV captures and 3 aligned aerial multispectral scenes. This supplement separately reports four additional official aligned scenes whose viewpoints are not part of the UAV-focused main scope. The retained-subset diagnostic for raw unaligned input is reported as an applicability check and should not be interpreted as the main raw full-scene comparison protocol.

Depth diagnostics follow the relative reference-depth protocol when the reference source passes validation. Since raw UAV reconstructions are not in a globally calibrated metric coordinate system, all thresholds are relative-depth ratios rather than metric distances. We exclude failed dense-reference or frame-misaligned attempts from quantitative tables.

Resolution-aligned comparison views. The dagger marker ([†]) denotes UMGS renders bilinearly resized to the external adapter resolution solely for common-mask comparison. These rows are comparison views, not distinct training variants, and are not used to define the UMGS method.

Grouped split sanitization. For official aligned scenes with inconsistent split and transform metadata, we sanitize splits at the capture-group level against the available transforms. A D/G/R/RE/NIR group is removed only when one or more required members are absent from the transform file, making it unusable as a complete multispectral comparison unit. This is a protocol repair for data completeness, not a performance-based selection.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM XXXX-XXXX/2026/5-ART
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

Table 1. Raw self_m3m retained-subset diagnostic. This is an applicability check for the external pipeline under its retained subset, not the main raw full-scene comparison.

Method	PSNR	SSIM	LPIPS	4-band RMSE	SAM ^o
UMGS-I (retained)	28.35	0.834	0.203	0.0502	3.98
UMGS-J (retained)	28.33	0.834	0.204	0.0503	4.00
MS-Splatting (default)	31.07	0.933	0.107	0.0379	3.95

Depth reference. Depth diagnostics are validation-gated reference-depth checks, not dense MVS or LiDAR ground truth. The current main evidence uses references that pass visual, support, and camera-frame checks. Attempts whose fused point clouds were empty, whose meshes collapsed, or whose rendered reference views were frame-misaligned are retained only as audit artifacts and are not averaged into the reported depth table. The geometry claims therefore concern relative reference-depth consistency under held-out probe views, not absolute dense geometry accuracy.

2 Additional Non-UAV Official Scene-Wise External Comparison

Table 2 reports the four additional official aligned scenes that are intentionally kept out of the main UAV-focused tables. Table 3 provides their scene-wise common-mask comparison between UMGS-J[†] and MS-Splatting. We also include the full seven-scene official aligned average for reproducibility.

3 Depth-Reference Status and Visualization

We keep the relative-depth evaluator and validation scripts in the released code, and report quantitative depth only for validated references. The main diagnostic table uses four validated scenes: self_m3m, raw001, raw003, and ms_golf. Other attempted references are retained as audit artifacts only. For self_m3m, we additionally provide 10-held-out visual grids over RGB/G/R/RE/NIR. These grids show that UMGS-I and UMGS-J produce smoother low-tail relative-error maps, while the retained MS-Splatting adapter has more local holes and patchy deviations despite stronger strict-threshold agreement on some pixels.

The retained raw diagnostic is intentionally separate from the full raw protocol. It evaluates the subset that the external pipeline can retain, not the complete raw UAV reconstruction path.

4 Why Unfrozen-Support Is Not a Product Model

The Unfrozen-support branch is included as a single-band upper-bound ablation. It starts from the same RGB anchor as UMGS, but allows band-specific geometry and opacity updates in Stage 2. This can improve selected per-band render metrics, but it breaks the shared-support contract required by model-level spectral products.

Table 2. Additional official aligned non-UAV scenes, reported separately from the UAV-focused main-paper suite. RGB metrics for UMGS-J use the shared UMGS-I RGB anchor. Spectral metrics average G/R/RE/NIR.

Method	RGB PSNR	RGB SSIM	MS PSNR	MS SSIM	MS LPIPS	Idx. RMSE	Idx. SSIM
UMGS-I	21.67	0.690	25.75	0.763	0.262	0.108	0.710
UMGS-J	21.67	0.690	25.75	0.763	0.262	0.108	0.710

Table 3. Additional non-UAV official aligned scene-level common-mask comparison between UMGS-J[†] and MS-Splatting. PSNR/SSIM/LPIPS average G/R/RE/NIR; RMSE and SAM are spectral 4-band metrics.

Scene	UMGS-J PSNR	MMS PSNR	UMGS-J SSIM	MMS SSIM	UMGS-J LPIPS	MMS LPIPS	UMGS-J RMSE	MMS RMSE	UMGS-J SAM	MMS SAM
ms_bud	24.77	25.09	0.797	0.781	0.225	0.257	0.0630	0.0601	6.46	7.34
ms_fruit	24.18	23.81	0.698	0.701	0.318	0.306	0.0642	0.0670	8.12	8.66
ms_garden	29.44	27.36	0.874	0.820	0.179	0.283	0.0378	0.0460	3.60	4.40
ms_tree	26.45	24.83	0.812	0.757	0.251	0.325	0.0525	0.0644	5.90	7.07

Table 4. Full seven-scene official aligned common-mask average, included only for reproducibility. Lower is better for LPIPS, 4-band RMSE, and SAM.

Method	PSNR	SSIM	LPIPS	4-band RMSE	SAM ^o
UMGS-J [†]	26.55	0.820	0.231	0.0544	5.75
MS-Splatting	26.03	0.803	0.272	0.0571	6.23

The failure is concrete in our audit logs. On ms_lake, Unfrozen-support improves several band PSNR values, but product construction fails with an xyz mismatch between bands:

$$\max |\Delta xyz| = 6.84 \times 10^{-2}.$$

Scratch-MS training diverges further and triggers Gaussian-count mismatch across bands. These variants can still produce independent band images, and one could compute post-hoc 2D ratios after rendering; however, those ratios are not defined over a common 3D Gaussian support and therefore are not used as valid model-level multispectral products in our protocol.

5 UMGS-J Initialization Ablation (UMGS-J vs UMGS-JR)

The main paper keeps UMGS-JR out of the main tables and reports UMGS-J as the representative joint branch. This section records the initialization ablation that motivated that decision. Both variants use the same geometry-locked joint multispectral training design and export per-band evaluation views from a unified checkpoint. They differ only in appearance-bank initialization.

UMGS-J initializes each spectral bank from the reset-appearance state. UMGS-JR initializes each spectral bank from the RGB-tied carrier.

Table 5 compares these variants on the two joint-branch pilot scenes used during method screening. UMGS-JR does not improve the branch: it degrades image/index quality on both pilot scenes.

6 Takeaway

The supplementary tables support three practical points used by the paper: (1) additional official aligned scenes remain available for reproducibility but are outside the main UAV-focused scope, (2) depth-reference evidence is reported only when the reference source passes validation, and (3) UMGS-JR is an initialization ablation rather than a promoted branch.

Table 5. UMGS-J and UMGS-JR quality/cost comparison from pilot runs. Band metrics are G/R/RE/NIR means. Index metrics are NDVI/GNDVI/NDRE means.

Scene	Variant	Band SSIM ↑	Band PSNR ↑	Band LPIPS ↓	Index MAE ↓	Index RMSE ↓	Index PSNR ↑	Index SSIM ↑	Time (s) ↓
self_m3m	UMGS-J	0.7657	25.330	0.2113	0.0607	0.0863	27.701	0.7943	8129
self_m3m	UMGS-JR	0.7434	24.966	0.2061	0.0623	0.0880	27.519	0.7854	7339
Vineyard A	UMGS-J	0.8804	28.690	0.1092	0.0339	0.0513	32.169	0.8948	5805
Vineyard A	UMGS-JR	0.8590	27.847	0.1266	0.0375	0.0574	31.326	0.8800	6730